# Deep Generative Models

Lecture 12: Energy-Based Models

Aditya Grover

UCLA

# Recap of last lecture



$x^{(j)} \sim P_{data}$
$j = 1, 2, \ldots, m$

$d(P_{data}, P_\theta)$
$P_{data}$
$P_\theta$
$\theta \in M$

- Energy-based models: $p_\theta(\mathbf{x}) = \frac{\exp\{f_\theta(\mathbf{x})\}}{Z(\theta)}$.
  - $Z(\theta)$ is intractable, so no access to likelihood.
  - Comparing the probability of two points is easy:
    $p_\theta(\mathbf{x}')/p_\theta(\mathbf{x}) = \exp(f_\theta(\mathbf{x}') - f_\theta(\mathbf{x}))$.
- Maximum likelihood training: $\max_\theta \{f_\theta(\mathbf{x}_{train}) - \log Z(\theta)\}$.
  - Contrastive divergence:

    $$\nabla_\theta f_\theta(\mathbf{x}_{train}) - \nabla_\theta \log Z(\theta) \approx \nabla_\theta f_\theta(\mathbf{x}_{train}) - \nabla_\theta f_\theta(\mathbf{x}_{sample}),$$

    where $\mathbf{x}_{sample} \sim p_\theta(\mathbf{x})$.

## Sampling from EBMs: MH-MCMC

Metropolis-Hastings Markov chain Monte Carlo (MCMC).

1. $\mathbf{x}^0 \sim \pi(\mathbf{x})$
2. Repeat for $t = 0, 1, 2, \cdots, T - 1$:
   - $\mathbf{x}' = \mathbf{x}^t + $ noise
   - $\mathbf{x}^{t+1} = \mathbf{x}'$ if $f_\theta(\mathbf{x}') \geq f_\theta(\mathbf{x}^t)$
   - If $f_\theta(\mathbf{x}') < f_\theta(\mathbf{x}^t)$, set $\mathbf{x}^{t+1} = \mathbf{x}'$ with probability
     $\exp\{f_\theta(\mathbf{x}') - f_\theta(\mathbf{x}^t)\}$, otherwise set $\mathbf{x}^{t+1} = \mathbf{x}^t$.

Properties:

- In theory, $\mathbf{x}^T$ converges to $p_\theta(\mathbf{x})$ when $T \to \infty$.

- In practice, need a large number of iterations and convergence
  slows down exponentially in dimensionality.

## Sampling from EBMs: unadjusted Langevin MCMC

Unadjusted Langevin MCMC:

1. $\mathbf{x}^0 \sim \pi(\mathbf{x})$
2. Repeat for $t = 0, 1, 2, \cdots, T - 1$:
   - $\mathbf{z}^t \sim \mathcal{N}(0, I)$
   - $\mathbf{x}^{t+1} = \mathbf{x}^t + \epsilon \nabla_{\mathbf{x}} \log p_\theta(\mathbf{x})|_{\mathbf{x}=\mathbf{x}^t} + \sqrt{2\epsilon}\mathbf{z}^t$

Properties:

- $\mathbf{x}^T$ converges to $p_\theta(\mathbf{x})$ when $T \to \infty$ and $\epsilon \to 0$.
- $\nabla_{\mathbf{x}} \log p_\theta(\mathbf{x}) = \nabla_{\mathbf{x}} f_\theta(\mathbf{x})$ for continuous energy-based models.
- Convergence slows down as dimensionality grows.

Sampling converges slowly in high dimensional spaces and is thus very expensive, yet we need sampling for **each training iteration** in contrastive divergence.

$x^{(j)} \sim P_{\text{data}}$
$j = 1, 2, \ldots, m$

$d(P_{\text{data}}, P_\theta)$
$P_{\text{data}}$
$P_\theta$
$\theta \in M$

**Goal**: Training without sampling

- Score Matching
- Noise Contrastive Estimation
- Adversarial training

**Energy-based model:** $p_\theta(\mathbf{x}) = \frac{\exp\{f_\theta(\mathbf{x})\}}{Z(\theta)}$

**(Stein) Score function:**

$$s_\theta(\mathbf{x}) := \nabla_\mathbf{x} \log p_\theta(\mathbf{x}) = \nabla_\mathbf{x} f_\theta(\mathbf{x}) - \underbrace{\nabla_\mathbf{x} \log Z(\theta)}_{0} = \nabla_\mathbf{x} f_\theta(\mathbf{x})$$

- Gaussian distribution
  $p_\theta(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$
  $\longrightarrow s_\theta(x) = -\frac{x-\mu}{\sigma^2}$

- Gamma distribution
  $p_\theta(x) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}$
  $\longrightarrow s_\theta(x) = \frac{\alpha-1}{x} - \beta$



$p_\theta(\mathbf{x})$ vs. $s_\theta(\mathbf{x})$

## Score matching

**Observation**

$s_\theta(\mathbf{x}) = \nabla_\mathbf{x} \log p_\theta(\mathbf{x})$ is independent of the partition function $Z(\theta)$.

Fisher divergence between $p(\mathbf{x})$ and $q(\mathbf{x})$:

$$D_F(p, q) := \frac{1}{2} E_{\mathbf{x} \sim p}[\|\nabla_\mathbf{x} \log p(\mathbf{x}) - \nabla_\mathbf{x} \log q(\mathbf{x})\|_2^2]$$

**Score matching:** minimizing the Fisher divergence between $p_\mathrm{data}(\mathbf{x})$ and the EBM $p_\theta(\mathbf{x})$

$$\frac{1}{2} E_{\mathbf{x} \sim p_\mathrm{data}}[\|\nabla_\mathbf{x} \log p_\mathrm{data}(\mathbf{x}) - s_\theta(\mathbf{x})\|_2^2]$$

$$= \frac{1}{2} E_{\mathbf{x} \sim p_\mathrm{data}}[\|\nabla_\mathbf{x} \log p_\mathrm{data}(\mathbf{x}) - \nabla_\mathbf{x} f_\theta(\mathbf{x})\|_2^2]$$

$$\frac{1}{2}E_{\mathbf{x}\sim p_{\text{data}}}[\|\nabla_{\mathbf{x}}\log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}}\log p_\theta(\mathbf{x})\|_2^2]$$

How to deal with $\nabla_{\mathbf{x}}\log p_{\text{data}}(\mathbf{x})$? Integration by parts!

$$\frac{1}{2}E_{x\sim p_{\text{data}}}[(\nabla_x\log p_{\text{data}}(x) - \nabla_x\log p_\theta(x))^2] \quad \text{(Univariate case)}$$

$$= \frac{1}{2}\int p_{\text{data}}(x)[(\nabla_x\log p_{\text{data}}(x) - \nabla_x\log p_\theta(x))^2]\mathrm{d}x$$

$$= \frac{1}{2}\int p_{\text{data}}(x)(\nabla_x\log p_{\text{data}}(x))^2\mathrm{d}x + \frac{1}{2}\int p_{\text{data}}(x)(\nabla_x\log p_\theta(x))^2\mathrm{d}x$$

$$- \int p_{\text{data}}(x)\nabla_x\log p_{\text{data}}(x)\nabla_x\log p_\theta(x)\mathrm{d}x$$

For the cross-correlation term:

$$- \int p_{\text{data}}(x)\nabla_x\log p_{\text{data}}(x)\nabla_x\log p_\theta(x)\mathrm{d}x$$

$$= - \int p_{\text{data}}(x)\frac{1}{p_{\text{data}}(x)}\nabla_x p_{\text{data}}(x)\nabla_x\log p_\theta(x)\mathrm{d}x$$

$$= \underbrace{-p_{\text{data}}(x)\nabla_x\log p_\theta(x)|_{x=-\infty}^{\infty}}_{=0} + \int p_{\text{data}}(x)\nabla_x^2\log p_\theta(x)\mathrm{d}x$$

$$= \int p_{\text{data}}(x)\nabla_x^2\log p_\theta(x)\mathrm{d}x$$

# Score matching

Univariate score matching

$$\frac{1}{2}E_{x \sim p_{\text{data}}}[(\nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_\theta(x))^2]$$

$$= \frac{1}{2}\int p_{\text{data}}(x)(\nabla_x \log p_{\text{data}}(x))^2 \mathrm{d}x + \frac{1}{2}\int p_{\text{data}}(x)(\nabla_x \log p_\theta(x))^2 \mathrm{d}x$$

$$- \int p_{\text{data}}(x)\nabla_x \log p_{\text{data}}(x)\nabla_x \log p_\theta(x)\mathrm{d}x$$

$$= \underbrace{\frac{1}{2}\int p_{\text{data}}(x)(\nabla_x \log p_{\text{data}}(x))^2 \mathrm{d}x}_{\text{const.}} + \frac{1}{2}\int p_{\text{data}}(x)(\nabla_x \log p_\theta(x))^2 \mathrm{d}x$$

$$+ \int p_{\text{data}}(x)\nabla_x^2 \log p_\theta(x)\mathrm{d}x$$

$$= E_{x \sim p_{\text{data}}}[\frac{1}{2}(\nabla_x \log p_\theta(x))^2 + \nabla_x^2 \log p_\theta(x)] + \text{const.}$$

Multivariate score matching

$$\frac{1}{2}E_{\mathbf{x} \sim p_{\text{data}}}[\|\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_\theta(\mathbf{x})\|_2^2]$$

$$= E_{\mathbf{x} \sim p_{\text{data}}}\left[\frac{1}{2}\|\nabla_{\mathbf{x}} \log p_\theta(\mathbf{x})\|_2^2 + \mathrm{tr}(\underbrace{\nabla_{\mathbf{x}}^2 \log p_\theta(\mathbf{x})}_{\text{Hessian of } \log p_\theta(\mathbf{x})})\right] + \text{const.}$$

## Score matching

1. Sample a mini-batch of datapoints
   $\{\mathbf{x}_1, \mathbf{x}_2, \cdots, \mathbf{x}_n\} \sim p_{\text{data}}(\mathbf{x})$.
2. Estimate the score matching loss with the empirical mean

$$\frac{1}{n} \sum_{i=1}^{n} \left[ \frac{1}{2} \|\nabla_{\mathbf{x}} \log p_\theta(\mathbf{x}_i)\|_2^2 + \text{tr}(\nabla_{\mathbf{x}}^2 \log p_\theta(\mathbf{x}_i)) \right]$$

$$= \frac{1}{n} \sum_{i=1}^{n} \left[ \frac{1}{2} \|\nabla_{\mathbf{x}} f_\theta(\mathbf{x}_i)\|_2^2 + \text{tr}(\nabla_{\mathbf{x}}^2 f_\theta(\mathbf{x}_i)) \right]$$

3. Stochastic gradient descent.
4. No need to sample from the EBM!

**Caveat**

Computing the trace of Hessian $\text{tr}(\nabla_{\mathbf{x}}^2 \log p_\theta(\mathbf{x}))$ is in general very expensive for large models.

Some solutions: Denoising score matching (Vincent 2010) and

# Recap



$$x^{(j)} \sim P_{\text{data}}$$
$$j = 1, 2, \ldots, m$$

$$d(P_{\text{data}}, P_\theta)$$

$$P_{\text{data}} \qquad P_\theta$$

$$\theta \in M$$

Distances used for training energy-based models.

- KL divergence = maximum likelihood.

$$\nabla_\theta f_\theta(\mathbf{x}_{data}) - f_\theta(\mathbf{x}_{sample}) \quad \text{(contrastive divergence)}$$

- Fisher divergence = score matching.

$$\frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}}[\|\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}} f_\theta(\mathbf{x})\|_2^2]$$

## Noise contrastive estimation

Learning an energy-based model by contrasting it with a noise distribution.

- Data distribution: $p_{\text{data}}(\mathbf{x})$.
- Noise distribution: $p_n(\mathbf{x})$. Should be analytically tractable and easy to sample from.
- Training a discriminator $D_\theta(\mathbf{x}) \in [0, 1]$ to distinguish between data samples and noise samples.

$$\max_\theta E_{\mathbf{x} \sim p_{\text{data}}}[\log D_\theta(\mathbf{x})] + E_{\mathbf{x} \sim p_n}[\log(1 - D_\theta(\mathbf{x}))]$$

- Optimal discriminator $D_{\theta^*}(\mathbf{x})$.

$$D_{\theta^*}(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_n(\mathbf{x})}$$

## Noise contrastive estimation

What if the discriminator is parameterized by

$$D_\theta(\mathbf{x}) = \frac{p_\theta(\mathbf{x})}{p_\theta(\mathbf{x}) + p_n(\mathbf{x})}$$

The optimal discriminator $D_{\theta^*}(\mathbf{x})$ satisfies

$$D_{\theta^*}(\mathbf{x}) = \frac{p_{\theta^*}(\mathbf{x})}{p_{\theta^*}(\mathbf{x}) + p_n(\mathbf{x})} = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_n(\mathbf{x})}$$

Equivalently,

$$p_{\theta^*}(\mathbf{x}) = \frac{p_n(\mathbf{x})D_{\theta^*}(\mathbf{x})}{1 - D_{\theta^*}(\mathbf{x})} = p_{\text{data}}(\mathbf{x})$$

**Noise contrastive estimation for training EBMs**

Energy-based model:

$$p_\theta(\mathbf{x}) = \frac{e^{f_\theta(\mathbf{x})}}{Z(\theta)}$$

The constraint $Z(\theta) = \int e^{f_\theta(\mathbf{x})} \mathrm{d}\mathbf{x}$ is hard to satisfy.

**Solution**: Modeling $Z(\theta)$ with an additional trainable parameter $Z$ that *disregards* the constraint $Z = \int e^{f_\theta(\mathbf{x})} \mathrm{d}\mathbf{x}$.

$$p_{\theta,Z}(\mathbf{x}) = \frac{e^{f_\theta(\mathbf{x})}}{Z}$$

With noise contrastive estimation, the optimal parameters $\theta^*, Z^*$ are

$$p_{\theta^*,Z^*}(\mathbf{x}) = \frac{e^{f_{\theta^*}(\mathbf{x})}}{Z^*} = p_{\text{data}}(\mathbf{x})$$

The optimal parameter $Z^*$ is the correct partition function, because

$$\int \frac{e^{f_{\theta^*}(\mathbf{x})}}{Z^*} \mathrm{d}\mathbf{x} = \int p_{\text{data}}(\mathbf{x}) \mathrm{d}\mathbf{x} = 1 \implies Z^* = \int e^{f_{\theta^*}(\mathbf{x})} \mathrm{d}\mathbf{x}$$

## Noise contrastive estimation for training EBMs

The discriminator $D_{\theta,Z}(\mathbf{x})$ for probabilistic model $p_{\theta,Z}(\mathbf{x})$ is

$$D_{\theta,Z}(\mathbf{x}) = \frac{\frac{e^{f_\theta(\mathbf{x})}}{Z}}{\frac{e^{f_\theta(\mathbf{x})}}{Z} + p_n(\mathbf{x})} = \frac{e^{f_\theta(\mathbf{x})}}{e^{f_\theta(\mathbf{x})} + p_n(\mathbf{x})Z}$$

Noise contrastive estimation training

$$\max_{\theta,Z} E_{\mathbf{x} \sim p_{\text{data}}}[\log D_{\theta,Z}(\mathbf{x})] + E_{\mathbf{x} \sim p_n}[\log(1 - D_{\theta,Z}(\mathbf{x}))]$$

Equivalently,

$$\max_{\theta,Z} E_{\mathbf{x} \sim p_{\text{data}}}[f_\theta(\mathbf{x}) - \log(e^{f_\theta(\mathbf{x})} + Zp_n(\mathbf{x}))]$$

$$+ E_{\mathbf{x} \sim p_n}[\log(Zp_n(\mathbf{x})) - \log(e^{f_\theta(\mathbf{x})} + Zp_n(\mathbf{x}))]$$

Log-sum-exp trick for numerical stability:

$$\log(e^{f_\theta(\mathbf{x})} + Zp_n(\mathbf{x})) = \log(e^{f_\theta(\mathbf{x})} + e^{\log Z + \log p_n(\mathbf{x})})$$

$$= \text{logsumexp}(f_\theta(\mathbf{x}), \log Z + \log p_n(\mathbf{x}))$$

## Noise contrastive estimation for training EBMs

1. Sample a mini-batch of datapoints $\mathbf{x}_1, \mathbf{x}_2, \cdots, \mathbf{x}_n \sim p_{\text{data}}(\mathbf{x})$.
2. Sample a mini-batch of noise samples $\mathbf{y}_1, \mathbf{y}_2, \cdots, \mathbf{y}_n \sim p_n(\mathbf{y})$.
3. Estimate the NCE loss.

$$
\frac{1}{n} \sum_{i=1}^{n} [f_\theta(\mathbf{x}_i) - \text{logsumexp}(f_\theta(\mathbf{x}_i), \log Z + \log p_n(\mathbf{x}_i))
$$
$$
+ \log Z + p_n(\mathbf{y}_i) - \text{logsumexp}(f_\theta(\mathbf{y}_i), \log Z + \log p_n(\mathbf{y}_i))]
$$

4. Stochastic gradient ascent.
5. No need to sample from the EBM!

## Comparing NCE and GAN

Similarities:

- Both involve training a discriminator to perform binary classification with a cross-entropy loss.
- Both are likelihood-free.

Differences:

- GAN requires adversarial training or minimax optimization for training, while NCE does not.
- NCE requires the likelihood of the noise distribution for training, while GAN only requires efficient sampling from the prior.
- NCE trains an energy-based model, while GAN trains a deterministic sample generator.

## Flow contrastive estimation (Gao et al. 2020)

**Observations**:

- We need to both evaluate the probability of $p_n(\mathbf{x})$, and sample from it efficiently.
- We hope to make the classification task as hard as possible, i.e., $p_n(\mathbf{x})$ should be close to $p_{\text{data}}(\mathbf{x})$ (but not exactly the same).

**Flow contrastive estimation**:

- Parameterize the noise as a normalizing flow $p_{n,\phi}(\mathbf{x})$.
- Parameterize the discriminator $D_{\theta, Z, \phi}(\mathbf{x})$ as

$$D_{\theta, Z, \phi}(\mathbf{x}) = \frac{\frac{e^{f_\theta(\mathbf{x})}}{Z}}{\frac{e^{f_\theta(\mathbf{x})}}{Z} + p_{n,\phi}(\mathbf{x})} = \frac{e^{f_\theta(\mathbf{x})}}{e^{f_\theta(\mathbf{x})} + p_{n,\phi}(\mathbf{x})Z}$$

- Train the flow model to minimize $D_{JS}(p_{\text{data}}, p_{n,\phi})$:

$$\min_{\phi} \max_{\theta, Z} E_{\mathbf{x} \sim p_{\text{data}}}[\log D_{\theta, Z, \phi}(\mathbf{x})] + E_{\mathbf{x} \sim p_{n,\phi}}[\log(1 - D_{\theta, Z, \phi}(\mathbf{x}))]$$

Samples from SVHN, CIFAR-10, and CelebA datasets.

Image source: Gao et al. 2020.

## Adversarial training for EBMs

Energy-based model:

$$p_\theta(\mathbf{x}) = \frac{e^{f_\theta(\mathbf{x})}}{Z(\theta)}$$

Upper bounding its log-likelihood with a variational distribution $q_\phi(\mathbf{x})$:

$$
\begin{aligned}
E_{\mathbf{x} \sim p_{\text{data}}}[\log p_\theta(\mathbf{x})] \qquad &= E_{\mathbf{x} \sim p_{\text{data}}}[f_\theta(\mathbf{x})] - \log Z(\theta) \\
&= E_{\mathbf{x} \sim p_{\text{data}}}[f_\theta(\mathbf{x})] - \log \int e^{f_\theta(\mathbf{x})} \mathrm{d}\mathbf{x} \\
&= E_{\mathbf{x} \sim p_{\text{data}}}[f_\theta(\mathbf{x})] - \log \int q_\phi(\mathbf{x}) \frac{e^{f_\theta(\mathbf{x})}}{q_\phi(\mathbf{x})} \mathrm{d}\mathbf{x} \\
&\leq E_{\mathbf{x} \sim p_{\text{data}}}[f_\theta(\mathbf{x})] - \int q_\phi(\mathbf{x}) \log \frac{e^{f_\theta(\mathbf{x})}}{q_\phi(\mathbf{x})} \mathrm{d}\mathbf{x} \\
&= E_{\mathbf{x} \sim p_{\text{data}}}[f_\theta(\mathbf{x})] - E_{\mathbf{x} \sim q_\phi}[f_\theta(\mathbf{x})] + H(q_\phi(\mathbf{x}))
\end{aligned}
$$

Adversarial training

$$\max_\theta \min_\phi E_{\mathbf{x} \sim p_{\text{data}}}[f_\theta(\mathbf{x})] - E_{\mathbf{x} \sim q_\phi}[f_\theta(\mathbf{x})] + H(q_\phi(\mathbf{x}))$$

## Conclusion

- Energy-based models are very flexible probabilistic models with intractable partition functions.
- Sampling is hard and typically requires iterative MCMC approaches.
- Computing the likelihood is hard.
- Comparing the likelihood/probability of two different points is tractable.
- Maximum likelihood training by contrastive divergence. Requires sampling for each training iteration.
- Sampling-free training: score matching.
- Sampling-free training: noise contrastive estimation. Additionally provides an estimate of the partition function.
- Sampling-free training: adversarial optimization.
- Reference: *How to Train Your Energy-Based Models* by Yang Song and Durk Kingma.